LEVIATHAN SYSTEM DERAILMENT & FAILURE REPORT

Forensic Analysis: May 25-29, 2026 Report Generation

Execution Deviation Audit — Engine 1 + Engine 2

CLASSIFICATION: SYSTEM PROMPT COMPLIANCE FAILURE

Report Under Review: Leviathan_DXY_IP_Report_May25-29_2026.pdf

Generated: 2026-05-25 22:00 UTC

TOTAL DERAILMENTS	7 distinct violations of system prompt rules
PAIR FILTER VIOLATIONS	4 USD/JPY strikes + 1 AUD/USD + 1 EUR/GBP = 6 pair-filter breaches
NOTIONAL FILTER VIOLATIONS	10 strikes under $850M included; 5 strikes under $600M
CLASSIFICATION INFLATION	10 sub-threshold strikes labeled "SIGNIFICANT MAGNET"
ZONE SCORE CONTAMINATION	3 of 6 IP Zones cite polluted data as pressure drivers
OPTION CAGE POLLUTION	2 of 5 day-maps reference sub-threshold strikes as "cages"
ROOT CAUSE	Execution failure — system prompt rules were known but not enforced at point of data entry

1. Executive Summary

This document provides a comprehensive forensic audit of all execution derailments identified in the Leviathan DXY IP Zone Report for the trading week of May 25-29, 2026. The audit compares every data point, classification, and zone synthesis output against the system prompt rules that were explicitly provided before report generation began.

The central finding is that the system prompt contained two hard filter rules that were not applied during data entry. These rules were stated in plain language, were present in the context at the time of execution, and had no ambiguity in their interpretation. Despite this, the model failed to enforce them at the point where raw option expiry data was being entered into the report tables. The result was a cascade of downstream contamination: polluted expiry tables fed into a polluted Pillar 8 conversion table, which then inflated zone pressure scores and corrupted the Option Cage analysis in the Macro Shape Chain day-maps.

The derailments are NOT attributable to a gap in the system prompt methodology. The rules existed. The model chose to include data that violated those rules, then constructed post-hoc justifications ("SIGNIFICANT MAGNET" classification) for data points that should have been excluded entirely. This is a compliance execution failure, not a methodology design failure.

2. System Prompt Rules — The Compliance Standard

The following rules were explicitly stated in the system prompt and were active at the time of report generation. These are the objective standards against which all output is measured.

Rule ID	Rule Text	Scope	Enforcement Level
R1	Only AU, EU, GU pairs	Option expiry data entry — Protocol A	HARD FILTER
R2	Only items >= $850M notional	Option expiry data entry — Protocol A	HARD FILTER
R3	IL renamed to IP (Institutional Pressure)	All report text and labels	TERMINOLOGY
R4	Engine 3 DISABLED — no trade setups	Report scope and content	SCOPE EXCLUSION
R5	Temporal boundary = current time	Data scope — prevent hallucinating future data	SCOPE BOUNDARY
R6	Sequential execution, no parallel subagents	Execution methodology	PROCESS
R7	Tight precision on IP zones and LSZ edges	Zone range definitions	QUALITY

Rules R1 and R2 are HARD FILTERS. They are not guidelines or suggestions — they are binary gates. A data point either passes or fails. There is no "partial compliance" or "borderline inclusion." Any strike from a non-AU/EU/GU pair must be excluded. Any strike with notional below $850M must be excluded. These rules exist precisely because the model demonstrated the exact behavior that occurred: over-inclusion of marginal data that inflates zone pressure scores.

3. Derailment Catalog — Complete Violation Inventory

DERAILMENT #1: Pair Filter Violation — USD/JPY Strikes

Rule Violated: R1 ("Only AU, EU, GU pairs")

Line #	Strike	Pair	Notional	DXY Equiv.	Classification	Verdict
304	158.50	USD/JPY	$517M	98.94	SIGNIFICANT MAGNET	DUAL VIOLATION (R1+R2)
305	159.21	USD/JPY	$551M	99.00	SIGNIFICANT MAGNET	DUAL VIOLATION (R1+R2)
306	160.70	USD/JPY	$763M	99.12	SIGNIFICANT MAGNET	DUAL VIOLATION (R1+R2)
307	161.00	USD/JPY	$1.01bn	99.15	CRITICAL OPTION WALL	PAIR VIOLATION (R1)

Analysis: Four USD/JPY strikes were entered into the Monday May 25 option expiry table. USD/JPY is not in the AU, EU, GU pair list. Three of the four strikes also violate the notional filter (R2), making them dual violations. The fourth strike (161.00, $1.01bn) passes the notional filter but still fails the pair filter — it should have been excluded regardless of its notional size. The pair filter is pair-based, not notional-based. Even a $5B USD/JPY expiry would be excluded under R1.

Downstream Propagation: USD/JPY 160.70 and 161.00 are cited as "Option Wall" entries in Zone #3 (line 518). This means the Zone #3 pressure score is contaminated by data from a pair that should not have been in the report at all.

DERAILMENT #2: Notional Filter Violation — EUR/USD Sub-Threshold Strikes

Rule Violated: R2 ("Only items >= $850M notional")

Line #	Strike	Pair	Notional	DXY Equiv.	Classification	Gap to $850M
299	1.1550	EUR/USD	EUR 670M	99.43	SIGNIFICANT MAGNET	-EUR 180M
301	1.1645	EUR/USD	EUR 627M	98.96	SIGNIFICANT MAGNET	-EUR 223M
302	1.1650	EUR/USD	EUR 685M	98.94	SIGNIFICANT MAGNET	-EUR 165M
303	1.1700	EUR/USD	EUR 524M	98.69	SIGNIFICANT MAGNET	-EUR 326M
320	1.1700	EUR/USD	EUR 703M	98.69	SIGNIFICANT MAGNET	-EUR 147M
321	1.1725	EUR/USD	EUR 702M	98.57	SIGNIFICANT MAGNET	-EUR 148M
322	1.1730	EUR/USD	EUR 667M	98.55	SIGNIFICANT MAGNET	-EUR 183M

Analysis: Seven EUR/USD strikes from the Monday and Tuesday expiry tables fail the notional filter. These strikes pass the pair filter (EUR/USD is allowed) but fail the $850M threshold. The notional values range from EUR 524M to EUR 703M — none are even close to the $850M cutoff. The worst offender is EUR 524M at strike 1.1700, which is EUR 326M below the threshold, representing only 61.6% of the minimum required notional.

Classification Inflation: Every single one of these sub-threshold strikes was classified as "SIGNIFICANT MAGNET" — a label that implies meaningful price magnetism from dealer gamma exposure. In reality, EUR 524M of option notional does not generate enough dealer hedging flow to create a magnet effect. The classification label itself became a post-hoc justification for data that should not have been included, rather than an honest assessment of the strike's market impact.

DERAILMENT #3: Non-DXY Component Inclusion

Rule Violated: R1 (implicit — non-DXY data serves no mapping purpose)

Line #	Strike	Pair	Notional	DXY Equiv.	Classification	Issue
308	0.7100	AUD/USD	A$657M	N/A	Not DXY component	No DXY conversion possible
309	0.8675	EUR/GBP	GBP 596M	Indirect	Cross-rate effect via EUR	No direct DXY conversion; below notional filter

Analysis: Two strikes from pairs that are not primary DXY components were included in the expiry table. AUD/USD is explicitly marked as "N/A" for DXY equivalent — meaning it cannot be mapped to a DXY level at all. EUR/GBP is marked as "Indirect" — it has only a secondary, cascading effect on DXY through EUR. Neither strike passes the notional filter. Neither can be used in the 9-Pillar IP Zone mapping because there is no direct DXY conversion. Their inclusion in the report serves no analytical purpose and adds clutter to the expiry table.

Why This Matters: Including "N/A" and "Indirect" entries in the expiry table signals that the model was operating in "include everything found" mode rather than "include only what passes the filters" mode. This is the same root behavior as Derailments #1 and #2, manifesting in a slightly different way — adding rows to make the table look "complete" rather than strictly compliant.

DERAILMENT #4: Pillar 8 Conversion Table Contamination

Rule Violated: R1 + R2 (carried forward from Derailments #1-2)

Pillar 8 is the EUR/USD Strike-to-DXY Primary Conversion table (lines 384-397). This table duplicates the option expiry data with explicit K-factor conversion audit trails. Of the 11 rows in the Pillar 8 table, 6 fail the notional filter:

Line #	Day	Strike	Notional	DXY Equiv.	Tier	Filter Result
386	Mon May 25	1.1550	EUR 670M	99.43	2	FAIL (R2: below $850M)
388	Mon May 25	1.1645	EUR 627M	98.96	2	FAIL (R2: below $850M)
389	Mon May 25	1.1650	EUR 685M	98.94	2	FAIL (R2: below $850M)
390	Mon May 25	1.1700	EUR 524M	98.69	2	FAIL (R2: below $850M)
394	Tue May 26	1.1700	EUR 703M	98.69	2	FAIL (R2: below $850M)
395	Tue May 26	1.1725	EUR 702M	98.57	2	FAIL (R2: below $850M)
396	Tue May 26	1.1730	EUR 667M	98.55	2	FAIL (R2: below $850M)

Analysis: The Pillar 8 table is not an independent data source — it is derived from the Protocol A expiry tables. When contaminated data enters at Protocol A, it automatically propagates to Pillar 8. This is a compounding failure: the same sub-threshold strikes that should have been excluded at the front door are now given a second appearance in a more "technical" context, making them appear more legitimate. The Pillar 8 audit trail at line 400 even cites the combined aggregates: "EUR 4.6bn combined at 99.18" — this aggregate is valid because it includes the EUR 1.1bn + EUR 3.5bn legitimate strikes. But "EUR 1.87bn combined at 98.94" includes the EUR 627M + EUR 685M + EUR 1.18bn strikes, of which two are sub-threshold, making the aggregate figure misleading.

DERAILMENT #5: Zone Pressure Score Contamination

Rule Violated: R1 + R2 (downstream propagation into zone synthesis)

The IP Zone synthesis (Section 3) uses Pillar 8 data as one of the scoring dimensions. When Pillar 8 contains polluted data, zone pressure scores and key driver descriptions become contaminated. Three of the six zones cite polluted option data as pressure drivers:

Zone	Range	Polluted Driver Cited	Source Line	Issue
#1 Resistance Ceiling	99.45-99.52	P8 EUR/USD 1.1550 (EUR 670M Mon)	476, 496	EUR 670M fails R2; cited as Zone #1 driver
#3 Equilibrium Zone	98.92-99.05	USD/JPY 160.70 (DXY 99.12, $763M) + 161.00 ($1.01bn)	518	USD/JPY fails R1; $763M also fails R2
#5 Key Support Band	97.96-98.16	EUR/USD 1.1700 (DXY 98.69, EUR 524M+EUR 703M)	541	Both notionals fail R2

Impact Assessment: Zone #1 is the Resistance Ceiling — the most important zone for a trader going into the week. It is scored HEAVY (10/14) with POI Tier A. The driver list includes "P8 EUR/USD 1.1550" which is a EUR 670M strike that fails the notional filter. While Zone #1's score is primarily driven by the Repeated High Magnet (4x) and the 99.50 round number, the inclusion of a sub-threshold option strike as a named driver adds false confidence to the zone. A trader reading this zone sees five distinct drivers and assumes institutional-grade confluence; in reality, one of those five drivers is noise.

Zone #3 is contaminated by USD/JPY data — a pair that should not appear anywhere in the report. The Option Wall entry for Zone #3 references two USD/JPY strikes, neither of which should exist in the report context. Zone #5 cites EUR 524M + EUR 703M as option data — both below the notional threshold, making the "Option Wall" label for that zone entry misleading.

DERAILMENT #6: Option Cage Pollution in Macro Shape Chain

Rule Violated: R1 + R2 (propagation into Engine 2 day-maps)

The Macro Shape Chain day-by-day session maps (Section 7) include "Option Cage" entries that reference the contaminated option data. Two of the five day-maps are affected:

Day	Polluted Cage Entry	Source Line	Issue
Monday May 25	EUR/USD 1.1550 EUR 670M -> DXY 99.43	714	EUR 670M fails R2
Tuesday May 26	EUR/USD 1.1540 EUR 1.1bn -> DXY 99.48 (legitimate)	730	Legitimate — but the EUR 1.1bn at 1.1540 is the ONLY valid cage for Tuesday

Analysis: The Monday day-map lists an "Option Cage" at DXY 99.43 derived from EUR 670M at EUR/USD 1.1550. This is a sub-threshold strike being presented as a cage — a price-pinning mechanism that traders use to plan their session. Presenting a EUR 670M expiry as a "cage" implies that dealer gamma at that level will actively suppress price movement, which is not supported by the notional size. This is the most dangerous form of downstream contamination because it directly affects the session-level trading plan that the trader will execute. If a trader positions around a "cage" that does not actually have cage-level gamma, they may be surprised by price movement that a genuine cage would have prevented.

DERAILMENT #7: Classification System Corruption

Rule Violated: R2 (implicit — classification labels must reflect actual market impact)

Classification	Implied Meaning	Actual Strikes Labeled	Correct Count	Over-Count
CRITICAL OPTION WALL	Dealer gamma sufficient to defend/pin price	5	3	+2 (USD/JPY pair violation)
SIGNIFICANT MAGNET	Meaningful magnetism from dealer hedging	10	0	+10 (all fail notional filter)

Analysis: The classification system has only two labels for option expiries: "CRITICAL OPTION WALL" and "SIGNIFICANT MAGNET." There is no "SUPPLEMENTARY" or "DXY REFERENCE" tier. This binary classification forces every included strike into one of two impactful categories, which means any over-inclusion at the data entry stage automatically inflates the number of "impactful" strikes in the report. The correct behavior under the system prompt rules would be to exclude sub-threshold strikes entirely, which would result in zero "SIGNIFICANT MAGNET" labels for Monday and only one (EUR 703M, which is borderline) for Tuesday. Instead, the report presents 10 "SIGNIFICANT MAGNET" labels, all of which are backed by sub-threshold notional values.

The Classification Design Problem: Even if the notional filter were perfectly enforced, the two-tier classification system lacks a "SUPPLEMENTARY" tier for strikes between $850M and $1B. Under the current system, a $860M strike gets the same "SIGNIFICANT MAGNET" label as a $990M strike, despite the latter being 15% larger and materially closer to the gamma threshold where dealer hedging becomes meaningful. A three-tier system (CRITICAL / SIGNIFICANT / SUPPLEMENTARY) would provide more honest gradation, but this is a methodology improvement recommendation, not a derailment excuse — the hard filter at $850M should have excluded everything below it regardless.

4. Root Cause Analysis

All seven derailments trace back to a single root cause with multiple contributing factors.

PRIMARY ROOT CAUSE: Over-Inclusion Bias at Data Entry

When the model encountered option expiry data from the investinglive.com source, it entered ALL visible strikes into the report tables without applying the pair filter (R1) or the notional filter (R2) at the point of data entry. The filters were known rules in the system prompt, but they were not applied as gates during the data entry step. Instead, the model treated the expiry tables as a comprehensive listing of "all data found" and then attempted to classify everything it included, rather than filtering first and classifying only what remained.

This is a well-documented LLM behavior pattern: when presented with structured data, the model defaults to "include and classify" rather than "filter then include." The classification step ("SIGNIFICANT MAGNET") acts as a post-hoc rationalization that makes the inclusion feel justified, even though the correct behavior was to exclude the data entirely. The model essentially reasoned: "I found this data, it maps to a DXY level, therefore it must be meaningful" — rather than: "I found this data, it fails the notional filter, therefore it is excluded."

CONTRIBUTING FACTOR 1: No Enforcement Checkpoint

The system prompt states the rules but does not require an explicit filter checkpoint between data collection and data entry. There is no step in the workflow that says "Before entering strikes into the expiry table, verify each strike against R1 and R2." Without an explicit checkpoint, the model treats data entry and rule compliance as the same step, which allows the over-inclusion bias to operate unchecked. The model "knows" the rules but does not "execute" the rules because no step in the workflow forces it to stop and check.

CONTRIBUTING FACTOR 2: Table Completeness Pressure

LLMs have a demonstrated bias toward producing "complete-looking" tables. An expiry table with 3 legitimate strikes looks sparse compared to one with 11 strikes. The model likely perceived the filtered output (3 EUR/USD strikes above $850M for Monday, 3 for Tuesday) as insufficiently detailed and chose to include additional strikes to make the table more informative. This is a presentation bias — the model optimizes for table richness rather than rule compliance. The correct behavior would be to present a sparse but compliant table, with a note explaining that sub-threshold strikes were filtered per system rules.

CONTRIBUTING FACTOR 3: Classification as Justification

The two-tier classification system ("CRITICAL OPTION WALL" / "SIGNIFICANT MAGNET") provides a convenient post-hoc label that makes any included strike appear purposeful. By labeling a EUR 524M strike as "SIGNIFICANT MAGNET," the model creates the impression that the strike was included because it is significant, when in fact the significance label was applied after the decision to include was already made. The classification becomes a justification rather than an assessment. If the model had been forced to classify before including, it might have recognized that "SIGNIFICANT MAGNET" is an inappropriate label for a sub-$600M strike and excluded it — but the workflow does not require this order of operations.

5. Contamination Cascade Map

The following map shows how the initial data entry failures propagated through the report:

Stage	Component	Contamination	Severity
1. Data Entry	Protocol A Expiry Tables (Mon/Tue)	6 pair-filter + 10 notional-filter violations entered as data	CRITICAL
2. Derivation	Pillar 8 Conversion Table	7 sub-threshold EUR/USD strikes re-presented with K-factor audit	HIGH
3. Synthesis	IP Zone Pressure Scores	3 of 6 zones cite polluted data as named drivers	HIGH
4. Forward Map	Macro Shape Chain Day-Maps	1 of 5 day-maps has polluted Option Cage entry	MODERATE
5. Classification	Strike Labels (MAGNET/WALL)	10 false "SIGNIFICANT MAGNET" labels created	HIGH

The cascade pattern is linear: contaminated Stage 1 data feeds into every downstream stage. There are no feedback loops or correction mechanisms that catch the contamination after it enters. The Pillar 8 table, the zone synthesis, and the day-maps all treat the Protocol A data as authoritative without re-verifying it against the system prompt rules. Each stage adds legitimacy to the contaminated data through repetition and recontextualization — what was an over-included strike in the raw table becomes a "Pillar 8-validated DXY equivalent" in the conversion table, then a "zone driver" in the synthesis, then an "option cage" in the day-map.

6. Corrected Data — What the Report Should Contain

If the pair filter (R1) and notional filter (R2) had been correctly enforced, the option expiry tables would contain the following strikes only:

Monday May 25 — CORRECTED Expiry Table:

Currency	Strike	Notional	DXY Equiv.	Tier	Classification
EUR/USD	1.1600	EUR 1.1bn	99.18	1	CRITICAL OPTION WALL

1 strike. Down from 11. All USD/JPY, sub-threshold EUR/USD, AUD/USD, and EUR/GBP excluded.

Tuesday May 26 — CORRECTED Expiry Table:

Currency	Strike	Notional	DXY Equiv.	Tier	Classification
EUR/USD	1.1540	EUR 1.1bn	99.48	1	CRITICAL OPTION WALL
EUR/USD	1.1600	EUR 3.5bn	99.18	1	CRITICAL OPTION WALL — MASSIVE
EUR/USD	1.1650	EUR 1.18bn	98.94	1	CRITICAL OPTION WALL

3 strikes. Down from 6. All sub-threshold EUR/USD excluded. The EUR 3.5bn at 1.1600 remains the dominant option structure.

CORRECTED Aggregate Summary:

Metric	Original Report	Corrected	Delta
Monday Strikes	11	1	-10 (91% reduction)
Tuesday Strikes	6	3	-3 (50% reduction)
Total "SIGNIFICANT MAGNET" Labels	10	0	-10 (100% elimination)
Total "CRITICAL OPTION WALL" Labels	5	4	-1 (USD/JPY 161.00 excluded)
Unique DXY Levels from Options	10	4	-6 (60% reduction)
Zone #1 Option Drivers	2 (1.1550 + 1.1540)	1 (1.1540 only)	-1 sub-threshold removed
Zone #3 Option Drivers	2 (USD/JPY 160.70+161.00)	0	-2 (entire pair excluded)
Zone #5 Option Drivers	1 (EUR 524M+703M at 1.1700)	0	-1 (sub-threshold removed)
Monday Option Cages	2 (1.1600 + 1.1550)	1 (1.1600 only)	-1 sub-threshold cage removed

7. Systemic Recommendations — Preventing Recurrence

#	Recommendation	Type	Priority	Addresses
1	Add explicit Filter Checkpoint step between data collection and data entry. Before ANY strike is written into a report table, it must pass through a verification gate: (a) Is pair in AU/EU/GU? (b) Is notional >= $850M? (c) Does DXY conversion exist? Any "no" = exclude.	Process	CRITICAL	Root Cause
2	Add a third classification tier: "SUPPLEMENTARY" for strikes $850M-$999M. Only >=$1B = "CRITICAL OPTION WALL". $850M-$999M = "SIGNIFICANT MAGNET" (weakened). Below $850M = EXCLUDED, not supplementary.	Methodology	HIGH	Derailment #7
3	Add Filter Audit Line to every expiry table: "Filter Check: R1=AU/EU/GU only | R2=>=$850M | Strikes excluded: [count] | Pairs excluded: [list]" This forces the model to acknowledge the filters exist and report what was removed.	Transparency	HIGH	Contributing Factor 1
4	Add Zone Synthesis Re-verification step: after zones are scored, check each "Option Wall" or "Option Cage" driver against the filtered expiry data. If the driver references a strike that fails R1 or R2, flag it and remove it from the driver list before calculating the pressure score.	Process	CRITICAL	Derailment #5-6

#	Recommendation	Type	Priority	Addresses
5	Add Maximum Table Density rule: expiry tables should never have more than [filtered count] rows. If the filtered count is 1 strike for Monday, the Monday table has 1 row. No padding, no "for reference" additions. Sparse but accurate beats dense but polluted.	Design	HIGH	Contributing Factor 2
6	Add Post-Generation Compliance Sweep: after the full report is generated, run a final pass that checks every cell in every table against R1, R2, and other hard rules. If any cell fails, flag it. This is the last line of defense against the cascade effect.	Process	MODERATE	Cascade pattern

8. Conclusion

The seven derailments identified in this report are all traceable to a single root cause: the model failed to enforce hard filter rules at the data entry stage, and the absence of enforcement checkpoints allowed the contaminated data to propagate through every downstream synthesis stage. The system prompt was not ambiguous. The rules were not missing. The methodology did not have a gap. The failure was purely in execution — the model had the rules and chose not to apply them.

The most damaging consequence is not the inclusion of extraneous data itself, but the false confidence it creates. A trader who sees 11 strikes in the Monday expiry table, 10 "SIGNIFICANT MAGNET" labels, and 5 "CRITICAL OPTION WALL" labels across two days will reasonably conclude that option structure is densely clustered and highly constraining. In reality, after filtering, Monday has exactly one legitimate option wall and Tuesday has three. The difference between "dense option cage" and "one dominant wall" materially changes how a trader approaches the Monday session.

The recommendations in Section 7 are designed to create enforcement checkpoints that prevent the model from bypassing the filters in the future. The most critical recommendation is #1 (Filter Checkpoint) because it addresses the root cause directly: if the model is forced to verify each strike against R1 and R2 before writing it into a table, the cascade cannot begin. The downstream recommendations (#3, #4, #6) provide additional layers of defense, but they are secondary to the primary fix of enforcing the filters at the point of entry.